Lesson 1.1: The four decisions before the first prompt

Claude Certified Associate – Foundations (CCAO-F) · Module 1: Platform & Model Foundations · Lesson 1 · 11:55

Description

Most people open a new chat and start typing. But the setup they skipped decides how good the answer can get. This lesson teaches the four decisions Anthropic's official prep course puts before every prompt: the entry point (where the work happens), the capability features (what to switch on), the model (which Claude does the thinking) and the context (what Claude gets to see). Each one comes with the single question to ask yourself. You then see the framework in Anthropic's own words, in the prep course and the official exam guide, and in the Help Center's descriptions of web search, Research, artifacts and the merged "one Claude" layout. Next, one real task for this course goes through all four decisions, out loud, in Claude Desktop with Sonnet 5.5. The result is a one-page Module 1 roadmap, checked against the exam guide. The lesson ends with one exam-style scenario, reasoned out loud.

Learning outcomes

By the end of this lesson you can:

1. Name the four decisions you make before the first prompt: entry point, features, model and context.
2. Explain why they set the ceiling on what a prompt can do, and where each one sits on the exam.
3. Run a real task through all four, out loud, and check the result.

Chapters

Time	Chapter
00:00	Preview and outcomes
00:33	Why it matters
01:16	The four decisions
02:50	Straight from the source
04:43	Hands-on: one real task
10:24	Exam angle
11:20	Recap and quick check

The four decisions, one question each

#	Decision	The question to ask	What the lesson names	Exam task statement	Deeper in
1	Entry point	Where should this work happen?	Chat (steer one piece of work yourself) · Cowork (hand off a bigger, multi-step job) · Code (software; for this exam, knowing when to hand off to a developer). On the merged layout you just ask	None names it (see below)	1.2
2	Capability features	What does this task need beyond a conversation?	A Project (work that keeps coming back) · Research (many searches and a cited report) · an artifact (something you'd put in front of someone). Switch on what the job needs, not everything	3.1	1.2
3	Model	How much thinking, how fast, how often?	Haiku (fastest) · Sonnet (in between; this course's default) · Opus (most capable of the three)	3.2, 3.3	1.3
4	Context	What does Claude need to see, and what should stay out?	Your messages, your files, a Project's knowledge, memory from past chats. It has a limit, and clutter crowds out what matters	3.4	1.4

The four decisions, their order and the phrase "quality ceiling" are Anthropic's (prep course). The one-line questions and the decision-to-task-statement mapping are this course's own.

Entry point and the exam guide. "Entry point" is the prep course's term. It doesn't appear anywhere in the exam guide, and neither does Cowork. The prep course pairs entry point with features ("Select the appropriate Claude entry point and feature set for a given professional task"), so this course files it beside task 3.1. That is our reading, not the guide's. In the hands-on, Claude noticed the same gap on its own and said so.

Your app may look different: the Chat | Cowork switch

On September 16, 2026, Anthropic began merging Claude Cowork and chat into one Claude. It reaches accounts in stages, starting with Pro and Max plans, and there's nothing to turn on. The Help Center's test: "If you're on a Pro or Max plan and your message box still shows "Chat" and "Cowork" options, you don't have it yet." This lesson was recorded on 2026-10-04 on an account that still showed the switch. On the

merged layout, decision one gets simpler: "you don't need to decide which option better suits your task before getting started. Ask Claude for what you need, and it can decide which tool to use." There's also no web search toggle, because Claude searches the web when it helps. The other three decisions stay the same.

Key terms

- **Quality ceiling:** the best a session can produce with its setup. A better prompt can't lift it past a wrong setup (Anthropic's phrase).
- **Entry point:** where the work happens: which Claude surface you start in, such as Chat, Cowork or Code.
- **Capability feature:** a tool you switch on or ask for on top of the conversation: Projects, Research, artifacts, web search, connectors, skills.
- **Model:** which Claude does the thinking (Haiku, Sonnet or Opus), traded against speed and usage.
- **Context:** everything Claude can see when it answers: your messages, attached files, a Project's knowledge, memory. It has a limit.

What happened in the hands-on

Claude Desktop 2.9939.4 on Linux, Chat, Sonnet 5.5 at Medium effort, recorded on 2026-10-04. The four decisions, made out loud before typing:

1. **Entry point: Chat.** Cowork adds a setup row (Project or folder, Manual, Output) for handing off a bigger job. A one-page roadmap is something to steer yourself.
2. **Features: an artifact, asked for in the prompt.** No Project (it's a one-off; repeated work would belong in a Project). No Research (the sources were already in hand). Web search stayed on, but the prompt held Claude to the two files.
3. **Model: Sonnet 5.5.** Quoting from two files is a simple, well-defined job.
4. **Context: two files and nothing else.** Memory was switched off for this one chat in the + menu, so the account's memory didn't change. Two files were attached: the official exam guide PDF and a short source note with this course's Module 1 lesson list and the prep course's wording.

The prompt was three short paragraphs. Claude worked for about 55 seconds, without searching the web, and opened an artifact titled "Module 1 roadmap: four decisions before the first prompt". It quoted task statements 3.1 to 3.4 word for word from the guide. In the entry-point row it wrote: "The Exam Guide has no task statement about choosing an entry point, so none is quoted here." The prompt had asked for a task statement for every decision. Claude found none for entry point and said so instead of inventing one. Its closing note flagged that Cowork is taught in Lesson 1.2 but isn't in the exam blueprint. The 3.4 quote was checked on screen against page 5 of the guide. The other quotes and all four lesson titles were checked off camera. Everything matched.

The exact prompt, the source note and the roadmap Claude produced are in this lesson's student kit.

The exam link

Domain 3, Product and Model Selection (12% of scored items), task statements, quoted word for word from the guide (the guide lists them as bullets; the numbers 3.1–3.4 are this course's labels):

- 3.1 "Select appropriate Claude product features (Projects, research mode, chat, artifacts)"
- 3.2 "Differentiate between Claude model types (Haiku, Sonnet, Opus)"
- 3.3 "Align model selection with task requirements (cost, speed, quality)"
- 3.4 "Understand and manage context limitations and memory considerations (when to restart, summarize, or persist)"

This lesson is the framework. Lessons 1.2 (features and entry points), 1.3 (models) and 1.4 (context and memory) take each decision deeper.

The lesson's practice scenario (original, not from the exam). Every Monday, an analyst writes a status update from the same two policy documents and format notes. Each week she opens a new chat, uploads both again and re-explains the format, and the updates still come out uneven. The most appropriate change is **B: a Project, with the documents as knowledge and the notes as instructions.** Ask which decision is broken. It isn't the model, so a more capable model (A) just costs more. Research (C) goes out to the web, but her sources are two fixed documents. One long chat for every update (D) just piles up context. The real problem is setup she repeats every week. A Project gives her a self-contained workspace with its own knowledge and instructions (the Help Center: "Projects allow you to create self-contained workspaces with their own chat histories and knowledge bases"), so each week builds on the last. In the prep course's words, these decisions "determine whether session builds upon prior work or requires constant re-setup" (quoted as printed).

Quick check

Before you move on, take this lesson's quick check. It's a few short questions and takes a couple of minutes. Module 1's quiz comes after Lesson 1.4. The exam-style question in this lesson, the quick checks, the module quizzes and the practice exams are all original, written for this course against the official blueprint. They are never copies of the real exam.

Sources (all checked 2026-10-04)

- **Anthropic's official CCAO-F prep course**, Course 1 of 8, *Claude Platform & Model Foundations*: hero line, "About this course" and the four learning objectives (quoted word for word).
- **Anthropic, *Claude Certified Associate – Foundations Exam Guide*.** Version 1.0, effective July 2026. §6, Domain 3, task statements 3.1–3.4 (p. 5). Official download: https://everpath-course-content.s3-accelerate.amazonaws.com/instructor/6nizmqk8tpzpfjvt6qmmav7rh/public/1783542847/Claude+Certified+Associate+%E2%80%93+Foun

- **Claude Help Center:**
  - *When should I use web search, extended thinking, and research?*: https://support.claude.com/en/articles/11095361-when-should-i-use-web-search-extended-thinking-and-research (web search "with one or two tool calls"; Research "five or more tool calls over 1-3 minutes", with citations)
  - *Use research on Claude*: https://support.claude.com/en/articles/11088861-use-research-on-claude
  - *What are artifacts and how do I use them?*: https://support.claude.com/en/articles/17153992-what-are-artifacts-and-how-do-i-use-them (the "put in front of someone" test)
  - *What are projects?*: https://support.claude.com/en/articles/9517075-what-are-projects ("self-contained workspaces with their own chat histories and knowledge bases")
  - *Claude Cowork and chat are one Claude*: https://support.claude.com/en/articles/16761823-claude-cowork-and-chat-are-one-claude (the merged layout and its test)
  - *Change the model, effort, and thinking settings*: https://support.claude.com/en/articles/8664678-change-the-model-effort-and-thinking-settings
  - *What interfaces can I use to access Claude?*: https://support.claude.com/en/articles/8114487-what-interfaces-can-i-use-to-access-claude
  - *How large is the context window on paid Claude plans?*: https://support.claude.com/en/articles/8606394-how-large-is-the-context-window-on-paid-claude-plans (tools and connectors use context; keep only what you need)
  - *Usage limit best practices*: https://support.claude.com/en/articles/9797557-usage-limit-best-practices (tools, model choice and effort affect usage)
  - *Use Claude's chat search and memory to build on previous context*: https://support.claude.com/en/articles/11817273-use-claude-s-chat-search-and-memory-to-build-on-previous-context
- **Anthropic blog, *Claude Cowork and chat are now one Claude*** (September 16, 2026): https://claude.com/blog/cowork-is-now-claude
- **The app itself:** Claude Desktop 2.9939.4 on Linux, Max plan, observed 2026-10-04. The model picker's own descriptions: Opus 5.5 "For complex work and everyday tasks", Sonnet 5.5 "Most efficient for simpler tasks", Haiku 4.5 "Fastest for quick answers". Menus and labels can change between versions.

What you see in the lesson is a real account. The model picker is framed on the Opus, Sonnet and Haiku rows, the models this course and the exam teach. Settings → Account and Settings → Memory were never opened. A file-path suggestion list that flashed up for a few frames in the file dialog was covered in the edit.