Lesson 3.1: Discernment: checking accuracy and completeness against the source and the brief

Claude Certified Associate – Foundations (CCAO-F) · Module 3: Evaluating & Validating Output · Lesson 1 · 14:39

Description

Claude's work often looks finished: clear headings, real numbers, a tidy length. Anthropic's AI Fluency Index found that this is exactly when people check less, at least in the conversation. This lesson teaches you to check any output against its two references. The source is what it has to be true to; that's accuracy. The brief is what you asked for; that's completeness. You see where the guidance comes from: the official exam guide, Anthropic's official prep course, two Claude Academy courses, the AI Fluency Index and the Claude Help Center. Then one real job runs in Claude Desktop. Claude writes a short reading on the AI Fluency Index for learners starting this module, and the whole check runs live. The brief becomes a checklist before reading. Claims are traced to the source with Ctrl+F, and every requirement is ticked. The check finds one real slip, a number about conversations written as a number about users. One precise follow-up fixes it, and only that change is re-checked. The lesson ends with one exam-style scenario, reasoned out loud, and the cue to this lesson's quick check.

Learning outcomes

By the end of this lesson you can:

1. Tell accuracy from completeness, and check each against the right reference.
2. Trace every claim to the source, and tick every requirement in the brief.
3. Turn what you find into one precise fix, and re-check only what changed.

Chapters

Time	Chapter
00:00	Preview and outcomes
00:37	Why it matters
01:43	The idea: two references, two questions
03:29	Straight from the source
04:57	Hands-on: checking a real summary
12:50	Exam angle
13:59	Recap and quick check

The idea

Two references. Every output answers to two things. The **source** is the material it has to be true to: a document, some data, your notes. The **brief** is what you asked for: the task, what to cover, the length, the rules.

Two questions, in opposite directions.

	Accuracy	Completeness
Direction	From the output back to the source	From the brief to the output
The question	Is everything that's there true to the source?	Is everything I asked for there?
What you check	Every claim: a number, a name, a date, a cause	Every requirement in the brief
Verdicts	Matches · Changed · Not in source	Covered · Partial · Missing

Complete doesn't mean answering everything. If the source doesn't say, a complete answer says so, instead of filling the gap.

Why you need both. A wrong number is on the page, so at least you can see it. A missing point isn't on the page at all, so nothing looks wrong. You only notice a gap if you bring a list, and the brief is that list.

"Changed" is the subtle verdict. The words are close, but the meaning shifted: a dropped qualifier, or percent where the source says percentage points.

The source-and-brief check (this course's method, built from Anthropic's guidance):

1. Turn the brief into a checklist, before you read.
2. Trace every claim to the source.
3. Tick every requirement.
4. Fix precisely, keep the rest, and re-check what changed.

One honest limit. A clean trace means the output matches *this* source. It doesn't prove the source itself is right. Checking the sources themselves comes in Lesson 3.4.

Straight from the source (what the lesson shows)

- **Exam guide, page 4:** "Exam items are written against these objectives." Domain 2, Output Evaluation and Validation (21% of scored items, the biggest domain), first bullet: "Evaluate Claude-generated outputs for accuracy and completeness". This course labels it task 2.1. The guide lists its task statements as unnumbered bullets, so 2.1 is our label, not the guide's. The other five Domain 2 bullets fill the rest of Module 3 (Lessons 3.2 to 3.6).
- **Anthropic's official prep course, Evaluating & Validating Claude's Output:** it's about standing "behind every deliverable you put your name on". Its About section: "critically evaluate output against sources and standards". Its first learning objective: "Evaluate Claude-generated output for accuracy and completeness".
- **Claude Academy, AI Fluency: Framework and foundations, A closer look at Discernment:** Product Discernment evaluates the outputs themselves, accuracy first; Process Discernment looks for "logical errors, attention gaps, or inappropriate reasoning".
- **The AI Fluency Index (Anthropic Education Report):** its advice for polished outputs: "pause and ask: is this accurate? Is anything missing?" In conversations where artifacts are made, users are less likely to identify missing context (-5.2pp) or check facts (-3.7pp), at least in the conversation.
- **Claude Academy, AI capabilities and limitations, Steerability:** the "letter-over-spirit" failure: "the instruction was honored but the intent wasn't". So tick what you meant, not just what you wrote.
- **Claude Help Center, *Claude is providing incorrect or misleading responses*:** "Claude can write things that might look correct but are very mistaken." And, for answers built from web pages: "Original websites may contain important context or details not included in Claude's synthesis." Accuracy first, then completeness, in one article.

Key terms

- **Source:** the material an output must be true to: the document, data or notes it was made from.
- **Brief:** what you asked for: the task, what to cover, the length and the rules. It doubles as your completeness checklist.
- **Claim:** one checkable statement in the output: a number, a name, a date, a cause.
- **Accuracy:** true to the source: the right facts, the right numbers, the same meaning.
- **Completeness:** everything the brief asked for is there, or the output says plainly that the source doesn't have it.

What happened in the hands-on

Claude Desktop on Linux, Chat view, Sonnet 5.5 at Medium, recorded on 2026-10-06 on a real account. The job is real: learners starting Module 3 need a short reading on Anthropic's AI Fluency Index. The source is real too: the text of the Index page, copied into a plain-text file on 2026-10-06, so Claude worked from exactly the text that was checked. The error the check found was not planted. The two generations play in real time; only still pauses where nothing happened on screen were trimmed.

Setup. Nothing here needed keeping, so this ran in an incognito chat (the ghost icon). In the + menu, Web search was ticked, and it was unticked: the source is the file, and no facts should come from anywhere else. The menu was reopened to make sure. Then the file was attached with **Add files or photos**, and Claude showed one tool step: reading the uploaded file.

Step 1: the brief, written as a checklist. Every line of the prompt is something to tick off later:

> Task: Summarize the attached report as a short reading for beginners starting Module 3 of my course. Cover: what was studied and how, how many people took part, the iteration finding, what went up and what went down when artifacts were made (with the numbers), the three tips, and two limitations the report states. Length: under 250 words. Use only the attached report. If it doesn't say, say so.

One honest note: the course author read the report before writing the brief, and knew it never says how many people took part. It counts conversations. A reader would ask, so the item stayed in. The last line gives Claude permission to admit a gap instead of filling it.

1. **The reply:** "Reading: The AI Fluency Index", with sections for what was studied and how, how many took part, iteration, what went up and down when artifacts were made (each number in percentage points), three tips and two limitations. It looks finished, and that's exactly the moment the report warns about. The reply was copied with the button under it, so its words could be counted off camera.

Step 2: trace every claim. All 19 claims were traced off camera, during a pause. Four were traced live, each with Ctrl+F in the source file:

2. **A plain number:** "85.7% of conversations involved iterating and refining". In the source: 1 match, the same number and the same thing counted. **Matches.**
3. **A number with its label:** "checking facts (-3.7pp)". In the source: 1 match, "check facts (-3.7pp)". Right number, right behavior, same unit (pp is short for percentage points). **Matches.**
4. **A claim about meaning:** the findings are correlational, so the report can't say whether one behavior causes another. "correlational": 2 matches, both in the report's limitations. **Matches.**
5. **The last line of tip 3:** "Only 30% of users do this." Searching "30% of" finds 1 match: "In only 30% of conversations do users tell Claude how they'd like it to interact with them." The report counts conversations, not users, and one person can have many conversations. Close words, shifted meaning: **Changed.**
6. **The full trace:** 19 claims in all: 18 match, 1 changed, 0 not in the source.

Step 3: tick every requirement, from the top.

7. **What was studied and how:** 11 behaviors, one week, a privacy-preserving tool. Covered.
8. **How many people took part** (the subtle one): the reply says the report counts 9,830 conversations, not people, and doesn't say how many individuals took part. In the source, "9,830" has 2 matches, both conversations. When the source doesn't say, a complete answer

says so. Covered.
9. **The rest:** the iteration finding, all seven artifact numbers in percentage points, the three tips (tip 3 holds the changed claim) and two limitations (the one-week sample and correlational findings): all covered. Length, counted off camera: 218 words, under 250. Only the report: no outside facts, and it said so where the report is silent. **8 of 8 covered.** Completeness passes; accuracy has one problem left.

Step 4: fix precisely, then re-check what changed. Not a new prompt, and not a vague "make it better". One follow-up that names the exact sentence, says what the report says, and keeps everything else (the "say what to keep" habit from Lesson 2.3):

> Keep everything else exactly as it is. Change only: in tip 3, "Only 30% of users do this." The report says 30% of conversations, not 30% of users.

10. **The new version:** tip 3 now ends "Only 30% of conversations include this." Searching "30% of conversations" finds 1 match, word for word, in the same sentence as before: **Matches.**
11. **Its neighbors:** tip 2 above it and the limitations below are unchanged. Off camera, the two versions were compared line by line, and that one sentence is the only change. The rest had already been traced, so it wasn't checked again.
12. **The final reading:** 19 claims traced, all 19 match; 8 requirements, all 8 covered. It's ready to put in front of learners, and the author knows exactly why.

The exam link

Domain 2, Output Evaluation and Validation (21% of scored items, the biggest domain). This lesson serves the bullet this course labels **2.1**: "Evaluate Claude-generated outputs for accuracy and completeness". The guide lists its task statements as unnumbered bullets; the number 2.1 is this course's label, not the guide's. The rest of Module 3 covers the other five Domain 2 bullets: hallucinations (Lesson 3.2), inconsistencies and bias (3.3), fact-checking (3.4), human review (3.5) and shaping the deliverable (3.6).

The lesson's practice scenario (original, written for this course, not from the exam). A project coordinator asks Claude to summarize a 20-page vendor contract for her manager, who asked for five things: price, term, renewal, termination and liability. The summary is clear, and every figure in it matches the contract. What's the most appropriate next step before she sends it? A: send it, because every figure matches the contract. B: tick the manager's five items against the summary first. C: ask Claude to confirm that the summary is complete. D: ask Claude for a longer and more detailed summary instead. The answer is **B**. Ask what's been checked: accuracy, against the contract, yes; completeness, not at all. A confuses accurate with complete. C asks Claude to grade its own work, and the exam guide's own sample question says self-reported confidence isn't a reliable accuracy signal. D makes it longer, not complete. B checks the brief item by item, and anything missing goes back to the contract.

Quick check

Before you move on, take this lesson's quick check: a few short questions, a couple of minutes. The exam-style question in this lesson, the quick checks, the module quizzes and the practice exams are all original, written for this course against the official blueprint. They are never copies of the real exam.

Sources (all checked 2026-10-06)

- **Anthropic, *Claude Certified Associate – Foundations Exam Guide*.** Version 1.0, effective July 2026. §3 Intended Audience (p. 2), §6 Domain 2 and its detailed objectives (p. 4), §7 How to Prepare (p. 6), §8 Sample 1 and its rationale. Official download: https://everpath-course-content.s3-accelerate.amazonaws.com/instructor/6nizmqk8tpzpfjvt6qmmav7rh/public/1783542847/Claude+Certified+Associate+%E2%80%93+Foun
- **Anthropic's official CCAO-F prep course,** Course 3, *Evaluating & Validating Claude's Output* (public course page: hero, About, learning objectives): https://anthropic-partners.skilljar.com/evaluating-validating-claudes-output
- **Claude Academy, AI Fluency: Framework and foundations,** *A closer look at Discernment* (Key takeaways): https://academy.claude.com/courses/ai-fluency-framework-foundations/a-closer-look-at-discernment
- **Anthropic Education Report, *The AI Fluency Index*** (Claude Academy tutorial, originally published February 23, 2026; "Developing your own AI fluency", the artifact findings and the limitations). Also the hands-on source: https://academy.claude.com/tutorials/the-ai-fluency-index
- **Claude Academy, AI capabilities and limitations,** *Steerability* (Key takeaways: letter-over-spirit): https://academy.claude.com/courses/ai-capabilities-and-limitations/steerability
- **Claude Help Center:**
  - *Claude is providing incorrect or misleading responses. What's going on?* (March 16, 2026): https://support.claude.com/en/articles/8525154-claude-is-providing-incorrect-or-misleading-responses-what-s-going-on
  - *Upload files to Claude* (July 23, 2026: TXT is a supported document type): https://support.claude.com/en/articles/8241126-uploading-files-to-claude
- **Background, not shown in the lesson:**
  - Claude Academy, AI capabilities and limitations, *Working Memory*: https://academy.claude.com/courses/ai-capabilities-and-limitations/working-memory
  - Claude Platform Docs, *Define success criteria and build evaluations*: https://platform.claude.com/docs/en/test-and-evaluate/develop-tests
  - Claude Platform Docs, *Reduce hallucinations* ("External knowledge restriction", the idea behind "Use only the attached report"): https://platform.claude.com/docs/en/test-and-evaluate/strengthen-guardrails/reduce-hallucinations
- **The app itself:** Claude Desktop on Linux, Chat view, Sonnet 5.5 Medium, observed 2026-10-06: incognito chats, the + menu (Add files or photos, Web search), the file chip, the tool step that read the file, the copy button and the status lines. Menus, labels and status lines

can change between versions.

"Two references, two questions", the claim and requirement verdicts and the four-step source-and-brief check are this course's own teaching tools, built from the sources above. They are not an Anthropic list. The hands-on material is real: Anthropic's own published report, a real job for this course, and the course author's own brief. Every claim and requirement verdict is the course author's, recorded during the take.

What you see in the lesson is a real account. The hands-on ran in an incognito chat, so it left no chat history and created no artifact. Settings, Account and Memory were never opened. The new-chat screen shows the account's own greeting and suggestion chips.